

# Machine Learning I

## Supervised learning framework

Souhaib Ben Taieb

University of Mons



# Table of contents

Components of supervised learning

Examples of learning models

A note on  $E_{\text{in}}$  and  $E_{\text{out}}$  of MLE

Training and test errors

Estimation of the out-of-sample error  $E_{\text{out}}$

Model selection

Model selection with cross-validation

# Table of contents

Components of supervised learning

Examples of learning models

A note on  $E_{\text{in}}$  and  $E_{\text{out}}$  of MLE

Training and test errors

Estimation of the out-of-sample error  $E_{\text{out}}$

Model selection

Model selection with cross-validation

# Input and output variables

The **input variables**<sup>1</sup> are typically denoted using the symbol  $X$ . If we observe  $p$  different variables, we write  $X = (X_1, X_2, \dots, X_p)$ . The inputs belong to an *input space*  $\mathcal{X}$ .

- ▶ Examples:  $\mathcal{X} \subseteq \mathbb{R}^p$  or  $\mathcal{X} = \{0, 1\}^p$ .

The **output variable**<sup>2</sup> is typically denoted using the symbol  $Y$ . The output belongs to an *output space*  $\mathcal{Y}$ .

- ▶ Regression:  $\mathcal{Y} \subseteq \mathbb{R}$
- ▶ Classification (with  $K$  categories):  $\mathcal{Y} = \{\mathcal{C}_1, \mathcal{C}_2, \dots, \mathcal{C}_K\}$ 
  - ▶ Binary classification ( $K = 2$ ):  $\mathcal{Y} = \{-1, 1\}$  or  $\mathcal{Y} = \{0, 1\}$

---

<sup>1</sup>also called *predictors*, *independent variables*, *features*, *variables* or just *inputs*.

<sup>2</sup>also called the *response* or *dependent variable*.

# Joint distribution

We assume  $(X, Y) \sim p_{X,Y}$  where  $p_{X,Y}$  is a fixed unknown distribution which can be factorized as

$$p_{X,Y}(x, y) = p_X(x)p_{Y|X}(y|x),$$


where

- ▶ the marginal distribution  $p_X$  models uncertainty in the sampling of the inputs.
- ▶ the conditional distribution  $p_{Y|X}$  describes a stochastic (non-deterministic) relation between inputs and output.

Equivalently, we have

$$X \sim p_X \text{ and } Y|X = x \sim p_{Y|X}(\cdot|x).$$

# Joint distribution



Source: <https://tinyurl.com/19bdt531>

# Optimal predictions

Define a **loss function**  $L : \mathcal{Y} \times \mathcal{Y} \rightarrow [0, \infty)$ . Given a prediction  $\hat{y} \in \mathcal{Y}$  and the true (observed) value  $y \in \mathcal{Y}$ ,  $L(y, \hat{y})$  measures how far  $\hat{y}$  is from  $y$ .

Examples include the *squared error loss*  $L(y, \hat{y}) = (y - \hat{y})^2$ , the *absolute error loss*  $L(y, \hat{y}) = |y - \hat{y}|$ , and the *zero-one loss*  $L(y, \hat{y}) = \mathbb{1}\{y \neq \hat{y}\}$ , where  $\mathbb{1}\{\cdot\}$  is the indicator function.

The **optimal prediction function** which minimizes the **expected error** (or **expected risk**) is given by

$$f = \operatorname{argmin}_{h: \mathcal{X} \rightarrow \mathcal{Y}} \mathbb{E}_{\mathbf{x}, y} [L(y, h(\mathbf{x}))].$$

In practice, we cannot compute  $f$  since we **do not know**  $p_{\mathbf{X}, Y}$ .


# The dataset

The **dataset**, also called *training set*, is a set of  $n$  input-output pairs

$$\mathcal{D} = \{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\} = \{(x_i, y_i)\}_{i=1}^n,$$

where the data points  $(x_i, y_i)$  are assumed to be i.i.d. realizations of  $p_{X,Y}$ .

Each pair, also called an *example* or a *data point*, belongs to the *data space*  $\mathcal{X} \times \mathcal{Y}$ .



- ▶ Left figure:  $\mathcal{X} \subseteq \mathbb{R}$  ( $p = 1$ ) and  $\mathcal{Y} \subseteq \mathbb{R}$
- ▶ Right figure:  $\mathcal{X} \subseteq \mathbb{R}^2$  ( $p = 2$ ) and  $\mathcal{Y} = \{R, G, B\}$



# The supervised learning problem

Let  $\mathcal{H}$  be a **hypothesis set**, i.e. a set of prediction function (hypotheses) under consideration. An example is the linear hypothesis set

$$\mathcal{H} = \{h(x) = \beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p : \beta_0, \beta_1, \dots, \beta_p \in \mathbb{R}\}.$$

Given a dataset  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ , the goal of supervised learning is to learn a **prediction function**  $h : \mathcal{X} \rightarrow \mathcal{Y}$  to map new/unseen examples with minimal **prediction error**.

We can compute the **in-sample error** or **training error**<sup>3</sup>

$$E_{\text{in}}(h) = \frac{1}{n} \sum_{i=1}^n L(y_i, h(x_i)),$$

and solve the following optimization problem:

$$g_{\mathcal{D}} = \operatorname{argmin}_{h \in \mathcal{H}} E_{\text{in}}(h),$$

---

<sup>3</sup>Also called the **empirical risk**.

# The supervised learning problem

Ideally, we would like to select the hypothesis  $h \in \mathcal{H}$  which minimizes the **out-of-sample error**

$$E_{\text{out}}(h) = \mathbb{E}_{\mathbf{x}, y}[L(y, h(\mathbf{x}))], \quad (1)$$

and compute

$$g^* = \operatorname{argmin}_{h \in \mathcal{H}} E_{\text{out}}(h).$$

In summary, there are three different prediction functions:

$$f = \operatorname{argmin}_{h: \mathcal{X} \rightarrow \mathcal{Y}} E_{\text{out}}(h),$$

$$g^* = \operatorname{argmin}_{h \in \mathcal{H}} E_{\text{out}}(h),$$

and

$$g_{\mathcal{D}} = \operatorname{argmin}_{h \in \mathcal{H}} E_{\text{in}}(h).$$

# Summary

The **dataset**  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$  is composed of  $n$  input-output pairs  $(x_i, y_i)$ , which are i.i.d. realizations from an **unknown joint distribution**  $p_{X,Y}$  where  $X \in \mathcal{X}$  and  $Y \in \mathcal{Y}$ .

The **loss function**  $L : \mathcal{Y} \times \mathcal{Y} \rightarrow [0, \infty)$  allows us to measure the error we incur in predicting  $\hat{y}$  in place of  $y$ .

The **hypothesis set**  $\mathcal{H}$  is a set of prediction function under consideration. Each hypothesis  $h \in \mathcal{H}$  has an **in-sample error**  $E_{\text{in}}(h)$ , computed on  $\mathcal{D}$ , and an **out-of-sample error**  $E_{\text{out}}(h)$  which depends on  $p_{X,Y}$ .

Given  $\mathcal{H}$  and using  $\mathcal{D}$ , the **learning algorithm**  $\mathcal{A}$  picks the best hypothesis  $g$  from  $\mathcal{H}$  according to the loss function  $L$ .

Together, the hypothesis set and the learning algorithm are referred to as the **learning model**.

# Table of contents

Components of supervised learning

**Examples of learning models**

A note on  $E_{\text{in}}$  and  $E_{\text{out}}$  of MLE

Training and test errors

Estimation of the out-of-sample error  $E_{\text{out}}$

Model selection

Model selection with cross-validation

# Linear models

Let us consider a regression problem where  $x \in \mathbb{R}^p$ .


- ▶ The **hypothesis set** for linear models is given by

$$\mathcal{H} = \{h(x) = \beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p : \beta_0, \beta_1, \dots, \beta_p \in \mathbb{R}\}$$

- ▶ One **learning algorithm** is the *(ordinary) least squares* method.

# Linear models

Example with  $p = 2$ .



# K-Nearest Neighbors (KNN) model



- ▶ **Input:**  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ , number of neighbors  $K \leq n$ , and new input  $x_*$
- ▶ **Classification** with  $M$  categories ( $y \in \{C_1, \dots, C_M\}$ )
  - ▶ Find the  $K$  nearest points to  $x_*$  in  $\mathcal{D}$ , denoted  $\mathcal{N}_*$ .

$$\hat{p}(C_m \mid x = x_*) \approx \frac{1}{K} \sum_{i \in \mathcal{N}_*} \mathbb{1}\{y_i = C_m\} \quad (m = 1, 2, \dots, M)$$

- ▶
- $$h(x_*) = C_{m^*} \text{ where } m^* = \operatorname{argmax}_m \hat{p}(y = C_m \mid x = x_*)$$
- ▶ **Regression**
  - ▶ Find the  $K$  nearest points to  $x_*$  in  $\mathcal{D}$ , denoted  $\mathcal{N}_*$ .

$$h(x_*) = \frac{1}{K} \sum_{i \in \mathcal{N}_*} y_i$$

# K-Nearest Neighbours (KNN) model



K-Nearest Neighbours (KNN) is one of the simplest machine learning model for both classification and regression.



# Parametric and non-parametric models

- ▶ In a **parametric model**, every hypothesis is uniquely defined by a **fixed number of parameters**.
- ▶ In a **non-parametric model**, we can not describe a hypothesis with a fixed number of parameters. Usually the number of “parameters” **grows with the size of the dataset**.
- ▶ Both parametric and non-parametric models have **hyper-parameters** (structural parameters), while parametric models also have **parameters**
  - ▶ For KNN,  $K$  is a **hyper-parameter**.
  - ▶ For linear models,  $p$  is a **hyper-parameter**, and the coefficients  $\beta_j$  are **parameters**.
- ▶ We also make a distinction between **linear** and **non-linear** models.

# Table of contents

Components of supervised learning

Examples of learning models

**A note on  $E_{\text{in}}$  and  $E_{\text{out}}$  of MLE**

Training and test errors

Estimation of the out-of-sample error  $E_{\text{out}}$

Model selection

Model selection with cross-validation

## $E_{\text{in}}$ for MLE

Recall that, given a set of possible distributions

$$\mathcal{H} = \{p(y; \theta) : \theta \in \Theta\},$$

and a dataset  $y_1, y_2, \dots, y_n$ , the maximum likelihood estimator is given by

$$\hat{\theta} = \operatorname{argmax}_{\theta \in \Theta} \sum_{i=1}^n \log p(y_i; \theta) = \operatorname{argmin}_{\theta \in \Theta} \underbrace{\frac{1}{n} \sum_{i=1}^n -\log p(y_i; \theta)}_{E_{\text{in}}(\theta)}. \quad (2)$$

Since each hypothesis  $h$  is completely characterized by  $\theta$ , this is equivalent to

$$g = \operatorname{argmin}_{h \in \mathcal{H}} E_{\text{in}}(h).$$

## $E_{\text{out}}$ for MLE

By (strong) law of large numbers,

$$\frac{1}{n} \sum_{i=1}^n -\log p(y_i; \boldsymbol{\theta}) \xrightarrow{n \rightarrow \infty} \mathbb{E}[-\log p(y; \boldsymbol{\theta})]$$

In other words, we can think of maximum likelihood estimation as trying to minimize

$$E_{\text{out}}(\boldsymbol{\theta}) = \mathbb{E}[-\log p(y; \boldsymbol{\theta})]$$

## $E_{\text{out}}$ for MLE

$$\operatorname{argmin}_{\theta \in \Theta} E_{\text{out}}(\theta) = \operatorname{argmin}_{\theta \in \Theta} \mathbb{E}[-\log p(y; \theta)] \quad (3)$$

$$= \operatorname{argmin}_{\theta \in \Theta} \mathbb{E}[\log p(y) - \log p(y; \theta)] \quad (4)$$

$$= \operatorname{argmin}_{\theta \in \Theta} \mathbb{E} \left[ \log \frac{p(y)}{p(y; \theta)} \right] \quad (5)$$

$$= \operatorname{argmin}_{\theta \in \Theta} \int \log \frac{p(y)}{p(y; \theta)} p(y) \, dy \quad (6)$$

$$= \operatorname{argmin}_{\theta \in \Theta} \text{KL}(p_{\theta}, p), \quad (7)$$

where  $\text{KL}(q, p)$  is the KL-divergence between two distributions  $q$  and  $p$ , which measures the discrepancy between the two distributions. Note that KL-divergence is not a distance measure (not symmetric).

# Table of contents

Components of supervised learning

Examples of learning models

A note on  $E_{\text{in}}$  and  $E_{\text{out}}$  of MLE

**Training and test errors**

Estimation of the out-of-sample error  $E_{\text{out}}$

Model selection

Model selection with cross-validation

## How does $E_{\text{out}}$ relates to $E_{\text{in}}$ ?

$$E_{\text{out}}(h) = E_{\text{in}}(h) + [E_{\text{out}}(h) - E_{\text{in}}(h)].$$

To obtain a small  $E_{\text{out}}(h)$ , we want

1. small  $E_{\text{in}}(h)$
2. small  $[E_{\text{out}}(h) - E_{\text{in}}(h)]$

Selecting the best hypothesis by minimizing  $E_{\text{in}}(h)$  only can be misleading.

# Training and test errors in regression

Let us consider multiple hypothesis sets  $\mathcal{H}_1, \mathcal{H}_2, \dots, \mathcal{H}_M$ , and define

$$g_m = \operatorname{argmin}_{h \in \mathcal{H}_m} E_{\text{in}}(h) = \operatorname{argmin}_{h \in \mathcal{H}_m} \frac{1}{n} \sum_{i=1}^n (y_i - h(x_i))^2, \quad (8)$$

for  $m = 1, 2, \dots, M$ . Here, we consider the squared error loss.

Let  $\mathcal{D}_{\text{test}} = \{(x'_i, y'_i)\}_{i=1}^{n'}$  be another sample (**independent** of  $\mathcal{D}$ ) where  $(x'_i, y'_i) \stackrel{\text{i.i.d.}}{\sim} p_{X,Y}$ . We define the **testing/test error** of a hypothesis  $h$  as



$$E_{\text{test}}(h) = \frac{1}{n'} \sum_{i=1}^{n'} (y'_i - h(x'_i))^2.$$

Note that the test set **is not involved** in the learning process, and is only used to evaluate the hypothesis  $h$ .

Let us compare  $E_{\text{in}}(g_m)$  and  $E_{\text{test}}(g_m)$  for  $m = 1, 2, \dots, M$ .





# Training and test errors in regression



Black: true curve  
Orange: linear regression  
Blue/green: nonlinear regression

Grey: Training MSE  
Red: Test MSE  
Dashed: Minimum test MSE



# Training and test errors in regression



Black: true curve  
Orange: linear regression  
Blue/green: nonlinear regression

Grey: Training MSE  
Red: Test MSE  
Dashed: Minimum test MSE

# Training and test errors in regression



Black: true curve

Orange: linear regression

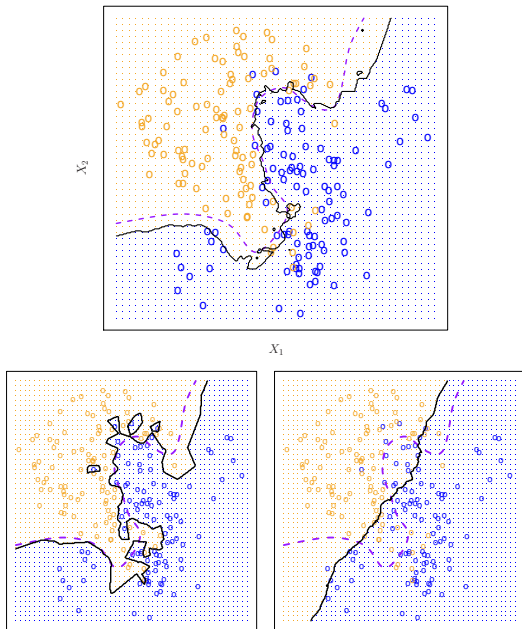
Blue/green: nonlinear regression

Grey: Training MSE


Red: Test MSE

Dashed: Minimum test MSE


# Training and test errors in classification



# Training and test errors in classification



# A fundamental picture



# Training and test errors

Consider

$$f = \operatorname{argmin}_{h: \mathcal{X} \rightarrow \mathcal{Y}} E_{\text{out}}(h),$$

$$g^* = \operatorname{argmin}_{h \in \mathcal{H}} E_{\text{out}}(h),$$

and

$$g = \operatorname{argmin}_{h \in \mathcal{H}} E_{\text{in}}(h).$$

How different are  $E_{\text{out}}(g)$  and  $E_{\text{out}}(f)$ ?

# The approximation-generalization tradeoff

The difference between the out-of-sample error of  $g$  and  $f$  can be decomposed as follows

$$E_{\text{out}}(g) - E_{\text{out}}(f) = \underbrace{[E_{\text{out}}(g^*) - E_{\text{out}}(f)]}_{\text{Approximation error}} + \underbrace{[E_{\text{out}}(g) - E_{\text{out}}(g^*)]}_{\text{Estimation error}}$$

- ▶ **Approximation error** is how far the entire hypothesis set is from  $f$ . Larger hypothesis sets have lower approximation error.
- ▶ **Estimation error** is how good  $g$  is with respect to the best in the hypothesis set. Larger hypothesis sets have higher estimation error because it is harder to find a good prediction function based on limited data.

This is called the **approximation-generalization** tradeoff.



# Table of contents

Components of supervised learning

Examples of learning models

A note on  $E_{\text{in}}$  and  $E_{\text{out}}$  of MLE

Training and test errors

**Estimation of the out-of-sample error  $E_{\text{out}}$**

Model selection

Model selection with cross-validation

# How to estimate the out-of-sample error?

The Holy Grail of machine learning is to find a good in-sample estimate of the out-of-sample error.

$$E_{\text{out}}(h) = E_{\text{in}}(h) + \underbrace{[E_{\text{out}}(h) - E_{\text{in}}(h)]}_{\text{overfit penalty}}, \quad h \in \mathcal{H}.$$

There are essentially two classes of methods:

1. **Directly estimate**  $E_{\text{out}}$  using resampling methods, i.e. by resampling the data set. Examples include **cross-validation** and **bootstrap** methods.
2. **Estimate the overfit penalty** and **add it** to  $E_{\text{in}}$ .



# Validation-set approach

- ▶ Randomly divide the dataset into two parts: a **training set** and a **validation set**.
- ▶ The model is fit on the training set, and the fitted model is used to predict the responses for the observations in the validation set



- ▶ The resulting **validation-set error** provides an estimate of the **out-of-sample error**.
- ▶ The idea of a **validation set** is almost identical to that of a **test set**. However, although the validation set will not be used for training, it will be used in making decisions during the learning process. A test set is not involved in any part of the learning process.


# Validation-set approach



*Left panel shows single split; right panel shows multiple splits*

# K-fold Cross-validation

- ▶ Divide the data set into  $K$  different parts.
- ▶ Remove one part, fit the model on the remaining  $K - 1$  parts, and compute the prediction error on the omitted part.
- ▶ Repeat  $K$  times taking out a different part each time
- ▶ By averaging the  $K$  prediction errors we obtain an estimate of the out-of-sample error, i.e. the prediction error for new observations (not used in training).



## **$K$ -fold Cross-validation**

Let the  $K$  parts be  $A_1, A_2, \dots, A_K$ , where  $A_k$  denotes the indices of the observations in part  $k$ . There are  $n_k$  observations in part  $k$ : if  $n$  is a multiple of  $K$ , then  $n_k = n/K$ .

Compute

$$E_{\text{cv}} = \sum_{k=1}^K \frac{n_k}{n} \text{Err}_k$$

where, for example,

$$\text{Err}_k = \sum_{i \in A_k} \frac{1}{n_k} (y_i - h^{(-k)}(x_i))^2 \quad (\text{squared error loss}),$$


or

$$\text{Err}_k = \sum_{i \in A_k} \frac{1}{n_k} \mathbb{1}\{y_i \neq h^{(-k)}(x_i)\} \quad (\text{zero-one loss}),$$

and  $h^{(-k)}(x_i)$  is the prediction for observation  $i$ , obtained from the hypothesis trained with data where part  $k$  is removed.

# Leave-one out cross-validation

Setting  $K = n$  yields  $n$ -fold or **leave-one out cross-validation** (LOOCV).







# **$K$ -fold Cross-validation - Bias and variance tradeoff**

- ▶ Each training set is only  $(K - 1)/K$  as big as the original data set. So the estimates of prediction error will be biased upwards.
- ▶ Bias minimized when  $K = n$  (LOOCV).
- ▶ But variance increases with  $K$  since the estimates from each fold are more correlated (as there are more overlapping observations in each part) and hence their average can have higher variance.
- ▶ Empirical observation:  $K = 5$  or  $K = 10$  provide a good compromise for this bias-variance tradeoff.



# 10-fold Cross-validation



*Left panel shows single split; right panel shows multiple splits*

# Table of contents

Components of supervised learning

Examples of learning models

A note on  $E_{\text{in}}$  and  $E_{\text{out}}$  of MLE

Training and test errors

Estimation of the out-of-sample error  $E_{\text{out}}$

**Model selection**

Model selection with cross-validation

# The model selection problem

We need a way of **assessing** and **selecting** the best model among multiple competing models.

- ▶ Which input variables should we choose?
- ▶ How many neighbours  $K$  should we use with KNN?
- ▶ ...
- ▶ More generally, how do we choose the **hyper-parameters** of a machine learning model?

# The model selection procedure

## 1. Model generation

Generate a set of **candidate model structures** (hypothesis set) among which the best one is to be selected. If applicable, estimate the parameters of each candidate model, i.e. fit the model by minimizing the **in-sample (training) error** (*parametric identification*).

## 2. Model assessment/validation

Evaluate/validate the model's performance by computing a **validation error**, i.e. an estimate of the out-of-sample error.

## 3. Model selection

Select the final model structure in the set that has been proposed by **model generation** and assessed by **model validation**. We typically select the model structure that minimizes the **validation error**.

Note that we **refit** the model on the entire dataset after model selection.

# Table of contents

Components of supervised learning

Examples of learning models

A note on  $E_{\text{in}}$  and  $E_{\text{out}}$  of MLE


Training and test errors

Estimation of the out-of-sample error  $E_{\text{out}}$

Model selection

**Model selection with cross-validation**

# The one standard error rule



Legend ?

Choose the simplest model whose CV error is no more than one standard error above the model with the lowest CV error

## Standard error of cross-validation estimate

It is very helpful to assign a quantitative notion of **variability** to the cross-validation error estimate. Assuming  $n_k = n/K$ , we can write

$$\text{Var}(E_{\text{cv}}) = \text{Var}\left(\sum_{k=1}^K \frac{1}{K} \text{Err}_k\right) \approx \frac{1}{K^2} \sum_{k=1}^K \text{Var}(\text{Err}_1) = \frac{1}{K} \text{Var}(\text{Err}_1)$$

- ▶ This is an approximation since  $\text{Err}_1, \dots, \text{Err}_K$  are not i.i.d.
- ▶ This approximation is valid for small  $K$  (e.g.,  $K = 5$  or  $10$ ) but not really for big  $K$  (e.g.,  $K = n$ ), because then the quantities  $\text{Err}_1, \dots, \text{Err}_K$  are highly correlated.
- ▶ We can compute the standard deviation or standard error of the cross-validation error estimate as

$$\frac{1}{\sqrt{K}} \text{sd}\{\text{Err}_1, \text{Err}_2, \dots, \text{Err}_K\},$$

where *sd* denotes the *empirical standard deviation*.

# Cross-validation: right and wrong


Consider a simple regression procedure applied to a dataset with 500 predictors and 50 samples:

1. Find the 5 predictors having the largest correlation with the response
2. Apply linear regression using only these 5 predictors

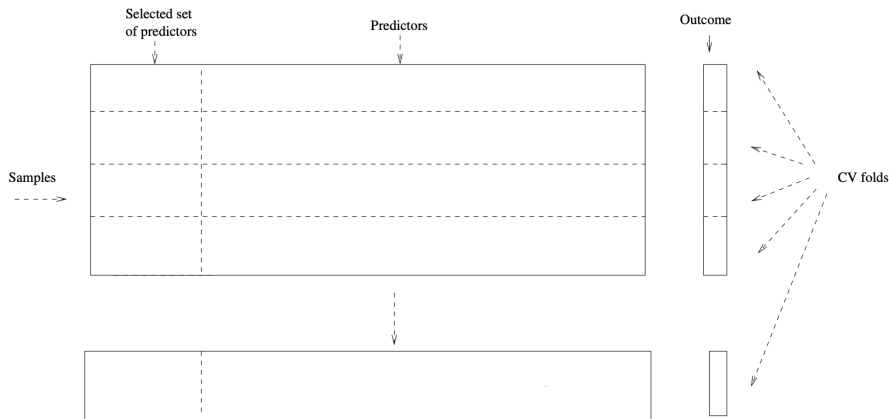
How to use cross-validation to estimate the out-of-sample error of this procedure?



# Cross-validation: wrong way



# Cross-validation: right way



→ Every aspect of the learning process that involves using the data — variable selection, scaling, etc — must be cross-validated.